

CS110 sp25 HW5

Due: 2025/5/7 23:59

Complete this homework either by writing neatly by hand or using [typst](#). You can find the .typ file on Piazza.

1 TRUE OR FALSE

Fill in your answer (T or F) in the table below.

1. When the same address is accessed multiple times using a cache, it primarily benefits from spatial locality.
2. If a cache system changes from direct-mapped to N-way set-associative ($N > 1$), the number of index bits in the address breakdown increases.
3. Higher cache associativity decreases hit time, miss rate, and miss penalty simultaneously.
4. Assume a direct-mapped cache with 8-byte blocks, 2 sets, using LRU replacement, and initially empty. Given 8-bit addresses, accessing addresses from 0x00 to 0xFF sequentially will result in no cache hits.

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 4 |
| F | F | F | F |

2 SET-ASSOCIATIVE CACHES

Assume a 12-bit address space. All cache lines are shown in the table below.

| Set index | Tag 1 | Tag 2 |
|-----------|-------|-------|
| 0 | | |
| 1 | | |
| 2 | | |
| 3 | | |

1. Assume loading a 16-byte struct at address 0x00F requires only 1 cache lookup, but loading one at 0x011 requires 2 lookups. Fill in the cache system parameters in the table below.

| | |
|---|-----------------------|
| Cache block size | 32 bytes |
| Total Capacity | 256 bytes |
| Level of set associativity | 2-way set associative |
| # of cache blocks | 8 |
| # of sets | 4 |
| Address Breakdown (Tag Index Offset) | 5, 2, 5 |

2. For the sequence of addresses below, indicate whether each access results in a hit, miss, or replacement. Assume the cache is initially empty.

| Address | Cache Access Result (Hit/Miss/Replacement) |
|---------|--|
| 0x3A8 | Miss |
| 0x1A6 | Miss |
| 0x04C | Miss |
| 0x5AD | Replacement |
| 0x3B9 | Replacement |
| 0x44F | Miss |
| 0x1B3 | Replacement |
| 0x055 | Hit |
| 0x241 | Replacement |
| 0x45E | Replacement |
| 0x1A1 | Hit |
| 0x1A5 | Hit |
| 0x642 | Replacement |
| 0x444 | Hit |

3. Calculate the cache hit rate for the sequence of memory accesses above. How can you improve the cache hit rate without increasing the cache **capacity**?

Solution: Change the cache associativity to fully associative.

3 AMAT

Consider a 3-level cache system with the following parameters, where the CPU runs at 4GHz:

| Cache Level | Hit Time | Local Miss Rate |
|-------------|-----------|-----------------|
| L1 | 2 cycles | 8% |
| L2 | 12 cycles | 35% |
| L3 | 44 cycles | 10% |
| Memory | 80ns | - |

1. Calculate the global miss rate for the 3-level cache.

Solution: $8\% \times 35\% \times 10\% = 0.28\%$

2. Calculate the average memory access time (AMAT) for the CPU.

Solution:

$$\text{AMAT} = 2 \text{ cycle} + 8\% \times (12 \text{ cycle} + 35\% \times (44 \text{ cycle} + 10\% \times (80 \text{ ns})))$$

Since the CPU runs at 4GHz, a cycle takes 0.25ns.

$$\text{AMAT} = 0.5 \text{ ns} + 0.08 \times (3 \text{ ns} + 0.35 \times (11 \text{ ns} + 0.10 \times (80 \text{ ns}))) = 1.272 \text{ ns}$$

3. Now, suppose the L1 cache is changed from 2-way set-associative to fully associative. This change reduces the L1 local miss rate to **6%** but increases the L1 hit time by **60%**. Calculate the new AMAT.

Solution:

$$\text{AMAT} = 0.8 \text{ ns} + 0.06 \times (3 \text{ ns} + 0.35 \times (11 \text{ ns} + 0.10 \times (80 \text{ ns}))) = 1.379 \text{ ns}$$

4 CODE ANALYSIS

Consider the following code:

```
struct body {
    float x, y, z, r;
};

struct body bodies[64];

// check whether two physics bodies overlap in 3D space
bool is_collide(struct body a, struct body b);

int check_collision() {
    int count = 0;
    for (int i = 0; i < 64; i++) {
        for (int j = i + 1; j < 64; j++) {
            if (is_collide(bodies[i], bodies[j])) {
                count++;
            }
        }
    }
    return count;
}
```

Note:

- Assume the cache parameters are: 128 bytes capacity, 32 bytes per block, 2-way set associative.
- Elements in the bodies array are aligned to the cache lines.
- `sizeof(float) == 4`
- You can ignore what `is_collide` exactly does and assume each body structure is loaded only **once** within the inner loop iteration (i.e., `bodies[i]` and `bodies[j]` are loaded into registers).
- The variables `i`, `j` and `count` are stored in registers.
- Instruction cache is not considered.
- Assume the cache is initially empty.

1. Calculate the hit rate for the above code.

Solution: The size of struct body is 16 bytes, which is exactly the size of a cache line.

$$\text{Total accesses} = \sum_{i=0}^{63} 2 \times (64 - i - 1) = 4032$$

1. For outer loop iterations where $i \leq 54$
 - Each time i reaches an even number (0, 2, 4, ..., 54), we'll have 1 miss
 - This gives us 28 misses from accessing body[i]
 - For the inner loop, we'll have a miss for every other element due to our 2-element-per-block alignment
 - Since i and $i+1$ loops have the same miss pattern, the total inner loop misses are 972
2. For outer loop iterations where $55 \leq i \leq 56$
 - Total 7 misses.
 - After this two iterations, all elements from body[57] to body[63] are in the cache.
3. For outer loop iterations where $i \geq 57$
 - Total 16 cache lines.
 - when $i \geq 57$, the last 16 elements of body are all in the cache and all accesses are hits.

$$\text{Hit Rate} = 1 - \frac{28 + 972 + 7}{4032} = \frac{3025}{4032} = 75.0248\%$$

2. How can the hit rate be improved by modifying only the code (without changing the cache configuration)? Briefly explain your solution.

Solution:

```
#define tile_size 8
int check_collision() {
    int count = 0;

    for (int ii = 0; ii < 64; ii += tile_size) {
        for (int jj = ii; jj < 64; jj += tile_size) {
            for (int i = ii; i < ii + tile_size && i < 64; i++) {
                for (int j = (i < jj ? jj : i + 1); j < jj + tile_size && j < 64; j++) {
                    if (is_collide(bodies[i], bodies[j])) {
                        count++;
                    }
                }
            }
        }
    }
    return count;
}
```

Use a smaller tile to make sure that the inner loop elements are always in the cache.